

OpenStreetMap Sample Project

Data Wrangling with MongoDB

Stephan Ketterer

Map Area: Las Vegas, Nevada, United States

https://s3.amazonaws.com/metro-extracts.mapzen.com/las-vegas_nevada.osm.bz2

1. Problems Encountered in the Map

Even though I created a snippet file of only the first 100 entries, it was quite hard for me to find a special pattern that needed a different cleanup than the one used in the last programming assignment. Since the chosen map was also a part of the United States, I figured I will most likely encounter the same problems, especially since Las Vegas is a well known city, with many sights and a lot of conferences. This should mean that there exists where solid open street map data on it. So I went with the established clean up procedure that was used in the Case Study in Lesson 6.

Subsequently I tackled the problem with the zip codes sometimes containing "NV". As shown in my provided .py file I simply replaced the part of the string in question.

The addresses were formatted as described in the programming assignment and different kind of attributes were saved properly for later access with mongodb queries.

Zip code problem:

```
zip=db2.col2.aggregate([{"$match":{"address.postcode":{"$exists":1}}}, {"$group":{"_id":"$address.postcode", "count":{"$sum":1}}}, {"$sort":{"count":-1}}])
```

```
[{u'count': 46, u'_id': u'89122'}, {u'count': 40, u'_id': u'89109'}, {u'count': 35, u'_id': u'89015'}, {u'count': 25, u'_id': u'89101'},
```

{u'count': 24, u'_id': u'89005'}, {u'count': 20, u'_id': u'89121'},
{u'count': 19, u'_id': u'89119'}, {u'count': 16, u'_id': u'89142'},
{u'count': 14, u'_id': u'89002'}, {u'count': 14, u'_id': u'89118'},
{u'count': 11, u'_id': u'89103'}, {u'count': 11, u'_id': u'89139'},
{u'count': 10, u'_id': u'89169'}, {u'count': 9, u'_id': u'89104'},
{u'count': 8, u'_id': u'89113'}, {u'count': 8, u'_id': u'89014'},
{u'count': 8, u'_id': u'89123'}, {u'count': 7, u'_id': u'89183'},
{u'count': 7, u'_id': u'89052'}, {u'count': 7, u'_id': u'89107'},
{u'count': 7, u'_id': u'89117'}, {u'count': 7, u'_id': u'89146'},
{u'count': 6, u'_id': u'89128'}, {u'count': 6, u'_id': u'89135'},
{u'count': 6, u'_id': u'89149'}, {u'count': 6, u'_id': u'89102'},
{u'count': 6, u'_id': u'89012'}, {u'count': 5, u'_id': u'89074'},
{u'count': 5, u'_id': u'89108'}, {u'count': 4, u'_id': u'89081'},
{u'count': 4, u'_id': u'89115'}, {u'count': 4, u'_id': u'89129'},
{u'count': 3, u'_id': u'89131'}, {u'count': 3, u'_id': u'89130'},
{u'count': 3, u'_id': u'NV 89109'}, {u'count': 3, u'_id': u'89144'},
{u'count': 3, u'_id': u'89145'}, {u'count': 2, u'_id': u'89040'},
{u'count': 2, u'_id': u'89120'}, {u'count': 2, u'_id': u'NV 89052'},
{u'count': 2, u'_id': u'89134'}, {u'count': 2, u'_id': u'NV 89119'},
{u'count': 2, u'_id': u'89147'}, {u'count': 2, u'_id': u'89106'},
{u'count': 1, u'_id': u'NV 89142'}, {u'count': 1, u'_id': u'89154'},
{u'count': 1, u'_id': u'89148'}, {u'count': 1, u'_id': u'NV 89129'},
{u'count': 1, u'_id': u'NV 89145'}, {u'count': 1, u'_id': u'NV
89124'}, {u'count': 1, u'_id': u'NV 89107'}, {u'count': 1, u'_id':
u'Nevada 89113'}, {u'count': 1, u'_id': u'NV 89014'}, {u'count':
1, u'_id': u'NV 89134'}, {u'count': 1, u'_id': u'89147-4111'},
{u'count': 1, u'_id': u'89011'}, {u'count': 1, u'_id': u'89025'},
{u'count': 1, u'_id': u'NV 89117'}, {u'count': 1, u'_id': u'89179'},
{u'count': 1, u'_id': u'NV 89123'}, {u'count': 1, u'_id': u'89178'},
{u'count': 1, u'_id': u'89030'}, {u'count': 1, u'_id': u'89191'},
{u'count': 1, u'_id': u'NV 89030'}, {u'count': 1, u'_id': u'89109-
1907'}, {u'count': 1, u'_id': u'89044'}, {u'count': 1, u'_id':
u'89032'}, {u'count': 1, u'_id': u'89156'}, {u'count': 1, u'_id':
u'89161'}, {u'count': 1, u'_id': u'NV 89191'}, {u'count': 1, u'_id':
u'NV 89031'}, {u'count': 1, u'_id': u'89105'}, {u'count': 1, u'_id':
u'NV 89101'}]

After cleaning up this problem and reuploading it to mongodb, the same query gave the desired result:

```
zipcleaned=db3.myc.aggregate([{"$match":{"address.postcode":{"$exists":1}}}, {"$group":{"_id":"$address.postcode", "count":{"$sum":1}}}, {"$sort":{"count":-1}}, {"$limit": 20}])
```

```
[{u'count': 46, u'_id': u'89122'}, {u'count': 40, u'_id': u'89109'}, {u'count': 35, u'_id': u'89015'}, {u'count': 25, u'_id': u'89101'}, {u'count': 24, u'_id': u'89005'}, {u'count': 20, u'_id': u'89121'}, {u'count': 19, u'_id': u'89119'}, {u'count': 16, u'_id': u'89142'}, {u'count': 14, u'_id': u'89002'}, {u'count': 14, u'_id': u'89118'}, {u'count': 11, u'_id': u'89139'}, {u'count': 11, u'_id': u'89103'}, {u'count': 10, u'_id': u'89169'}, {u'count': 9, u'_id': u'89104'}, {u'count': 8, u'_id': u'89113'}, {u'count': 8, u'_id': u'89123'}, {u'count': 8, u'_id': u'89014'}, {u'count': 7, u'_id': u'89052'}, {u'count': 7, u'_id': u'89183'}, {u'count': 7, u'_id': u'89107'}]
```

As predicted, the faulty "NV" has been successfully cleaned from the data.

City names:

Looking at the city names:

```
cities=db2.col2.aggregate([{"$match":{"address.city":{"$exists":1}}}, {"$group":{"_id":"$address.city", "count":{"$sum":1}}}, {"$sort":{"count":-1}}])
```

```
[{u'count': 248, u'_id': u'Las Vegas'}, {u'count': 79, u'_id': u'Henderson'}, {u'count': 5, u'_id': u'North Las Vegas'}, {u'count': 5, u'_id': u'Boulder City'}, {u'count': 2, u'_id': u'Spring Valley'}, {u'count': 2, u'_id': u'las vegas'}, {u'count': 2, u'_id': u'Boulder City NV'}, {u'count': 2, u'_id': u'Las vegas'}, {u'count': 2, u'_id': u'Boulder City, NV'}, {u'count': 2, u'_id': u'Las Vagas'},
```

```
{u'count': 2, u'_id': u'Overton'}, {u'count': 2, u'_id': u'Las Vegas, NV'}, {u'count': 1, u'_id': u'Nellis AFB'}, {u'count': 1, u'_id': u'Moapa'}, {u'count': 1, u'_id': u'Whitney'}]
```

It is observable that the data is still not perfectly formatted. There are some misspellings (Las Vagas), but also just different capitalization that leads to different values for the same city. Also looking at "Boulder City, NV" it seems that when submitting data points, it is not always clear where to actually put the state abbreviation in the data.

Cleaning the "NV" occurrence was cleaned the same way as for the zip code. After applying it the query came out the following:

```
[{u'count': 248, u'_id': u'Las Vegas'}, {u'count': 79, u'_id': u'Henderson'}, {u'count': 5, u'_id': u'North Las Vegas'}, {u'count': 5, u'_id': u'Boulder City'}, {u'count': 2, u'_id': u'Overton'}, {u'count': 2, u'_id': u'Spring Valley'}, {u'count': 2, u'_id': u'Boulder City, '}, {u'count': 2, u'_id': u'Boulder City '}, {u'count': 2, u'_id': u'las vegas'}, {u'count': 2, u'_id': u'Las Vegas, '}, {u'count': 2, u'_id': u'Las vegas'}, {u'count': 2, u'_id': u'Las Vagas'}, {u'count': 1, u'_id': u'Nellis AFB'}, {u'count': 1, u'_id': u'Moapa'}, {u'count': 1, u'_id': u'Whitney'}]
```

Still not perfectly clean, but definitely a step in the right direction.

Missing and not properly formatted data on amenities:

When asking for the top 10 types of restaurants:

```
db2.c3.aggregate([{"$match":{"amenity":{"$exists":1}, "amenity":"restaurant"}}, {"$group":{"_id":"$cuisine", "count":{"$sum":1}}}, {"$sort":{"count":-1}}, {"$limit":10}])
```

The result is the following:

```
[{'u'count': 105, 'u'_id': None}, {'u'count': 16, 'u'_id': u'pizza'},  
{u'count': 15, 'u'_id': u'mexican'}, {u'count': 10, 'u'_id':  
u'american'}, {u'count': 9, 'u'_id': u'burger'}, {u'count': 6, 'u'_id':  
u'japanese'}, {u'count': 6, 'u'_id': u'chinese'}, {u'count': 5,  
u'_id': u'italian'}, {u'count': 4, 'u'_id': u'thai'}, {u'count': 3,  
u'_id': u'sushi'}]
```

It is obvious that this subset of data is far from being clean and properly formatted. The biggest number has no id which means the user gets no further information on what kind of restaurant he or she should expect. Digging in deeper, it is observable that some entries are clearly problematic. For example, out of the 3 sushi restaurants, there is a good chance that some or all of them are japanese. Also the tag "american" seems not to be well chosen, because that could mean almost any kind of cuisine. In summary, there are definitely problems here with the data, but

they cannot be solved with cleaning. What needs to be done here is an additional input and revisiting of the data itself.

2.Data Overview

File sizes

lasvegas.osm	175 MB
--------------	--------

lasvegas.json	190 MB
---------------	--------

Number of documents

```
db2.col2.find().count()
```

885656

Number of nodes

```
db2.col2.find({"type":"node"}).count()
```

799544

#Number of ways

```
db2.col2.find({"type":"way"}).count()
```

86112

#Number of unique users

```
db2.col2.distinct("created.user").__len__()
```

681

#Top contributing user

```
db2.col2.aggregate([{"$group":{"_id": "$created.user",  
"count":{"$sum":1}}}, {"$sort":{"count":-1}}, {"$limit":1}])  
{u'count': 254460, u'_id': u'alimamo'}
```

3. Additional Ideas

Contributor statistics:

The amount each user contributed varied greatly.

Top user contribution percentage ('alimamo')=28.73%

Top 10 users contribution = 74.33%

Additional data exploration using MongoDB queries:

Top 10 appearing amenities

```

db2.col2.aggregate([{"$match": {"amenity": {"$exists": 1}}},
                    {"$group": {"_id": "$amenity", "count":
{"$sum": 1}}},
                    {"$sort": {"count": -1}},
                    {"$limit": 10}
[{'u'count': 797, 'u'_id': 'u'parking'},
{'u'count': 531, 'u'_id': 'u'school'},
{'u'count': 364, 'u'_id': 'u'place_of_worship'},
{'u'count': 269, 'u'_id': 'u'fountain'},
{'u'count': 218, 'u'_id': 'u'restaurant'},
{'u'count': 181, 'u'_id': 'u'fast_food'},
{'u'count': 143, 'u'_id': 'u'fuel'},
{'u'count': 69, 'u'_id': 'u'fire_station'},
{'u'count': 67, 'u'_id': 'u'hospital'},
{'u'count': 63, 'u'_id': 'u'post_office'}]

```

Most popular cuisines

```

db2.col2.aggregate([{"$match":{"amenity":{"$exists":1},
"amenity":"restaurant"}}, {"$group":{"_id":"$cuisine",
"count":{"$sum":1}}}, {"$sort":{"count":-1}}, {"$limit":3}])
{'u'count': 105, 'u'_id': None},
{'u'count': 16, 'u'_id': 'u'pizza'},
{'u'count': 15, 'u'_id': 'u'mexican'}

```

#Places of worship


```
db2.col2.aggregate([{"$match":{"amenity":{"$exists":1},
"amenity":"place_of_worship"}},
{"$group":{"_id":"$religion", "count":{"$sum":1}}},
{"$sort":{"count":-1}}, {"$limit":10}])
```

```
[{u'count': 339, u'_id': u'christian'},
 {u'count': 10, u'_id': None},
 {u'count': 4, u'_id': u'jewish'},
 {u'count': 3, u'_id': u'bahai'},
 {u'count': 2, u'_id': u'muslim'},
 {u'count': 2, u'_id': u'buddhist'},
 {u'count': 1, u'_id': u'scientologist'},
 {u'count': 1, u'_id': u'sikh'},
 {u'count': 1, u'_id': u'unitarian_universalist'}, {
u'count': 1, u'_id': u'hindu'}]
```

#Number of Casinos

```
casinos=db2.col2.aggregate([{"$match":{"amenity":'casino'}},{"$group": {"_id":"casinos", "count": {"$sum":1}}}]])
```

```
[{u'count': 30, u'_id': u'casinos'}]
```

#Number of Banks

```
banks=db2.col2.aggregate([{"$match":{"amenity":"'bank'"}}, {"$group": {"_id":"null", "count": {"$sum":1}}}]])
```

```
[{u'count': 53, u'_id': u'banks'}]
```

Conclusion

Clearly the data for Las Vegas is quite extensive and someone can get a lot of information out of it. But at second glance, it is obvious that a lot of information is either missing, or not properly formatted yet. Taking a look at the top ten list of amenities, restaurant and fast_food are listed separately, where clearly there should be some connection or a common tag. The phenomenon of missing data can be observed when looking at houses of worship. Number two on the list is again identified as "None", even though just by looking at the overall numbers, one can speculate that most of them are probably of Christian denomination. I think a possible approach to fill out the missing information and getting a better map file would be cross referencing the database with as many applicable databases as possible. I think vast improvements could be made, for looking for common tags and names and enriching the osm database with that new data. Different companies must have different databases covering the same area, for example a utility company might have very detailed and specific data relevant to their needs that could be integrated into the osm database. Lastly I think the huge number of casinos in Las Vegas must have something to do with the in my opinion very large number of banks for a city of that size.

